
From a 1.57 m KNN baseline to a 0.75 m cascade

Standard protocol 0.75 m · strict nested-CV 0.94 m

Four required sections, told as one engineering story

Problem · Method · Dataset & splits · Results

- 題目定義 / Problem: one variable-length scan to (x, y) in metres, no GPS
- 做法 / Method: Set Transformer to grid classification to coarse-gated cascade
- 資料集切分 / Dataset: 1,812 scans, 4 splits (A main, C cross-time hardest)
- 實驗結果 / Results: 1.57 m to 0.75 m, validated by nested cross-validation

SECTION 1 / 4

題目定義 / Problem Definition

What exactly are we predicting, and why is it hard?

The task: turn one variable-length WiFi scan into a point on the map

Input = a set of (BSSID, RSSI); output = (x, y) in metres

- Input: one scan = a variable-length set of (BSSID, RSSI) pairs
- Output: robot position (x, y) on the 189.5 m² floor plan
- No GPS indoors — WiFi fingerprinting is the classic fallback
- Four difficulties: variable set, ambiguous RSSI, sparse coverage, 0.3 m label floor

Anatomy of one scan: a ragged bag of (BSSID, RSSI)

Same place, two scans, different lengths — order means nothing, presence means everything

- One scan = an unordered bag like {(ap_3f, -52), (ap_1a, -67), (ap_b2, -78), ...}
- Length varies scan-to-scan (~12 to ~40 APs); same spot need not list the same APs
- Force it into an 80-D vector and a typical scan is ~2/3 -100 padding (27.7 of 80)
- Order carries no meaning; only which APs are present and how strong they are

SECTION 2 / 4

資料集切分 / Dataset & Splits

1,812 real scans, a coverage hole, and four generalisation stress-tests


Every scan is a real robot run: ESP32 on a TurtleBot3 doing SLAM

Hardware-collected data, not simulation — and that dictates every dataset quirk

- ESP32-S3 sniffer rides a TurtleBot3 running Cartographer SLAM + AMCL
- Each scan ~ 3 s; ground truth (x, y) is AMCL's pose (~ 0.3 m noise)
- Coverage = wherever we drove — path-shaped, not a uniform grid
- Two sessions (morning + evening) to expose later cross-time drift

1,812 fingerprints in a 189.5 m² lab — but the upper region is a coverage hole

- 1,812 scans = morning 912 + evening 900; train/test split 1449/363
- 115 BSSIDs seen, 80 kept (≥ 10 hits); the rest are transient noise
- Ground truth from ROS AMCL, ~ 0.3 m noise — our hard error floor
- Coverage follows the driven path; the upper region is barely visited





Four splits, each stress-testing a different generalisation axis

A_random is the main number; C (cross-time) is the hardest

- A_random (MAIN): 80/20 over all data, test=363 — accuracy ceiling
- B_morning: same-session — isolates model quality from time drift
- C_morning to evening: train AM, test PM — the hardest, cross-time test
- D_stratified: early/late mix — guards against trajectory-order leakage

Coverage, not noise, is the real bottleneck

- Coverage, not noise, is the bottleneck — the thesis of this whole deck
- Within-cell RSSI std ~ 3.5 dBm; morning-to-evening drift only ~ 0.8 dBm
- Error spikes when few known APs match, then flattens
- Diagnosis: fix coverage, not the network — motivates GP-synth later



SECTION 3 / 4


做法 / Method

The climb: each method change fixes one named problem

Classic KNN fingerprinting sets the bar to beat at 1.57 m

- Scan to 80-D RSSI vector; unheard AP padded to -100 dBm (not heard)
- $k=5$ (a-priori, not tuned on test) nearest neighbours, distance-weighted pose
- Zero training, pure lookup: the reference point every net must beat
- Median error 1.57 m on Split A test


KNN $k=5$ (baseline, median 1.568 m)



Plain neural nets barely beat KNN — and reveal two lessons

- Lesson 1: a dense -100-padded vector wastes most of the capacity
- Lesson 2: regression averages multimodal answers into the void
- MLP regression: 1.30 m — only a sliver below KNN
- MaskedMDN: 1.37 m — a mixture head alone does not fix multimodality

MaskedMDN (NN baseline, median 1.20 m)




Why predict a distribution, not a single point

One scan can map to two plausible places — a point estimate splits the difference

- Symmetric layouts make one RSSI fingerprint match two real positions
- L2 regression's optimum is the mean — it lands in the empty middle
- MDN expresses multiple peaks but its readout still averages them
- Cleanest fix: classify over map cells, where peaks stay separate


Treating the scan as a SET, not a vector, unlocks 1.09 m

- Each AP = one token: BSSID embedding + scaled RSSI
- Set Transformer: 3 self-attention blocks, pooling-by-attention to 192-d
- Permutation-invariant (order doesn't matter); only real APs enter
- Median 1.09 m, -16% vs the MLP




How a Set Transformer reads a variable-length set

- Each (BSSID, RSSI) becomes a token: learned ID embedding + scaled RSSI
- Self-attention = APs compare to each other, building joint context
- Permutation-invariant: no positional encoding, order carries no meaning
- PMA: a learned query attends over all tokens to a single 192-d vector




GP-kriging synthetic data: the single biggest jump, -17% to 0.91 m

- Per-BSSID Gaussian Process learns a smooth position-to-RSSI field
- Sample 5000 synthetic scans from the empty upper region
- Split A: 1.09 to 0.91 m (-17%) — the project's largest single gain
- In-session win; cross-time drift is revisited at the end




How GP-kriging fabricates believable coverage

- One 2-D GP per AP maps position to a smooth RSSI field with uncertainty
- A KNN detection model decides whether each AP is even visible here
- Both models are fit on TRAINING positions only — never on test poses
- Sanity check: synth mean 27.1 APs/scan vs real 27.7 — well-matched



The training recipe: augment, 5 seeds, geometric-median ensemble

- Train on real + 5000 synthetic scans together, not in stages
- Augment: RSSI jitter ± 4 dBm and 10% random AP dropout per scan
- Train 5 seeds; combine with a geometric-median ensemble (outlier-robust)
- Same recipe across every neural milestone — a fair, fixed comparison




Why geometric median? It is outlier-robust: one bad seed cannot drag the ensemble, unlike a plain mean. Augmentation injects the sensor noise the model must tolerate.

Classifying on a masked grid beats regressing a coordinate: 0.88 m


- Predict $P(\text{cell})$ over a 40×33 grid (~ 0.4 m cells), not a coordinate
- Free-cell mask zeroes walls; centroid reads off the position
- Classification can express multiple peaks; regression cannot
- Median 0.91 to 0.88 m (5-seed ensemble)

Set Transformer + Heatmap (median 0.883 m, $\times 5$ -ens)



Our champion: a coarse 10x9 gate over the fine 40x33 grid lands 0.79 m


- Shared encoder feeds a coarse 10x9 + a fine 40x33 head
- Gate: $P_{\text{fine}}(\text{cell}) \times P_{\text{coarse}}(\text{parent region})$ before the centroid
- 0.79 m median (5-seed) — best of every config we tried
- Loss-weight tuning (mse_w 0.55) to 0.752 m — the reported headline



Why it works: the coarse gate kills the fine grid's false second peak

- Fine head alone is bimodal: correct peak + a plausible decoy
- Probability-weighted centroid lands between them, off target
- Coarse head is confident about the region, suppresses the decoy
- After gating: one clean peak, prediction snaps back on target

Why the cascade wins: the coarse gate removes fine-grid ambiguity (this sample: 2.59 m $\rightarrow$ 0.90 m)




SECTION 4 / 4

實驗結果 / Experimental Results

The full climb, the validation, the failures, and the live demo


The full climb: every method change bought real metres, 1.57 to 0.75 m

- Two biggest jumps: the set representation and synthetic coverage
- Right inductive bias, not raw capacity, drove every gain
- Green = kept milestone, red = abandoned variant
- Caption: set-transformer onward = 5-seed; KNN/MLP/MDN = single model




Across all four splits: in-session is strong, cross-time is the weak spot

- A_random: 0.75 m (MEASURED, main headline)
- B_morning: ~0.85 m (same-session, estimate)
- D_stratified: ~0.90 m (all-session deployment, estimate)
- C_morning to evening: ~2.3 m (cross-time, ESTIMATE) — the open problem



The whole distribution improved, not just the median



- Cascade CDF sits left of every other model at every percentile
- 27% of predictions within the 0.3 m AMCL floor
- Median 0.79 m, mean 1.12 m, p90 2.56 m
- >2 m tail from sparse regions + symmetric-layout ambiguity



Where to trust it: reliability tracks data density, not network size

- Per-cell mean error mirrors the sample-density map almost exactly
- Dense, well-travelled regions: well under 0.5 m
- Sparse upper region: error climbs above 2 m
- The fix is more data, not a bigger model


Where can you trust the model? Error is low where data is dense, high in the sparse upper region.



How we measured: leak-free nested 5-fold cross-validation

- Outer 5 folds: each fold is a held-out test, never touched during training
- Per fold we retrain everything from scratch on the other 4/5 — no shared weights
- Model/hparam selection happens INSIDE each fold, never on the outer test
- Result: 0.94 ± 0.04 m — the most pessimistic, most trustworthy number

Leak-free nested 5-fold CV: every fold is retrained from scratch





Each fold: retrain the whole model on its 4 training folds + regenerate GP-synth from those folds only, then predict the held-out fold. Stack all 5 -> 1,449 honest out-of-fold predictions. No sample ever scores a model that saw it. -> honest 0.94 m

0.650 m was a test-set artifact; 0.752 m is the reported headline


- Greedy ensemble of 33 candidates hit 0.650 m — but picked ON the test set
- Held-out val exposed it at 0.710 m — a ~ 0.07 m selection bias, so we dropped it
- Strictest leak-free nested 5-fold CV ($4/5$ data per fold) = 0.94 ± 0.04 m
- SWA (top literature method) moved $0.025 < \text{bootstrap SE } 0.037 \rightarrow$ not adopted

Split A median: 0.752 m standard · 0.94 m strict nested-CV · 0.650 m was test-tuned




Bigger and fancier consistently lost: at 1.4k samples, bias beats capacity

- Diffusion head 1.82 m: sampling noise on a deterministic task
- Cascade-big 1.7M params 0.93 m: capacity just overfits at this scale
- CNN floor-plan cross-attn 0.91 m: extra machinery, no extra signal
- (also: C-Mixup 0.94, A+B/CNN+cascade combo 0.89, 3-level+12k synth 0.80)




The error splits cleanly: 27% solved at the label floor, 16% unsolvable

- 27% of predictions sit at the 0.3 m AMCL floor (label-limited, not improvable)
- ~16% are >2 m outliers from WiFi symmetric ambiguity (two places, same RSSI)
- Worst-10 cluster in the sparse upper region and mirror-symmetric spots
- No model can separate identical fingerprints — only new signals can



Live in the lab: ESP32 to 9 ms CPU inference to live heatmap

- Real ESP32 streaming live scans; CPU-only inference ~ 9 ms/scan (no GPU)
- Live position + probability heatmap rendered on the floor plan
- 80-95% of live APs match the trained vocabulary (no drift)
- Precision standing still < 0.1 m spread; accuracy vs ground truth 0.4-0.8 m



Verdict: ~0.75 m is real; sub-0.65 m needs new data, not a deeper net

- Halved KNN error to a 0.752 m median, validated on real hardware
- Won by inductive bias (set input, grid classification, coarse gate), not capacity
- Validation matters: a 0.650 m test-tuned result didn't hold; nested-CV gives 0.94 m
- To break 0.6x: upper-region data, IMU/heading, temporal fusion — not depth

